

Engenharia de Software II

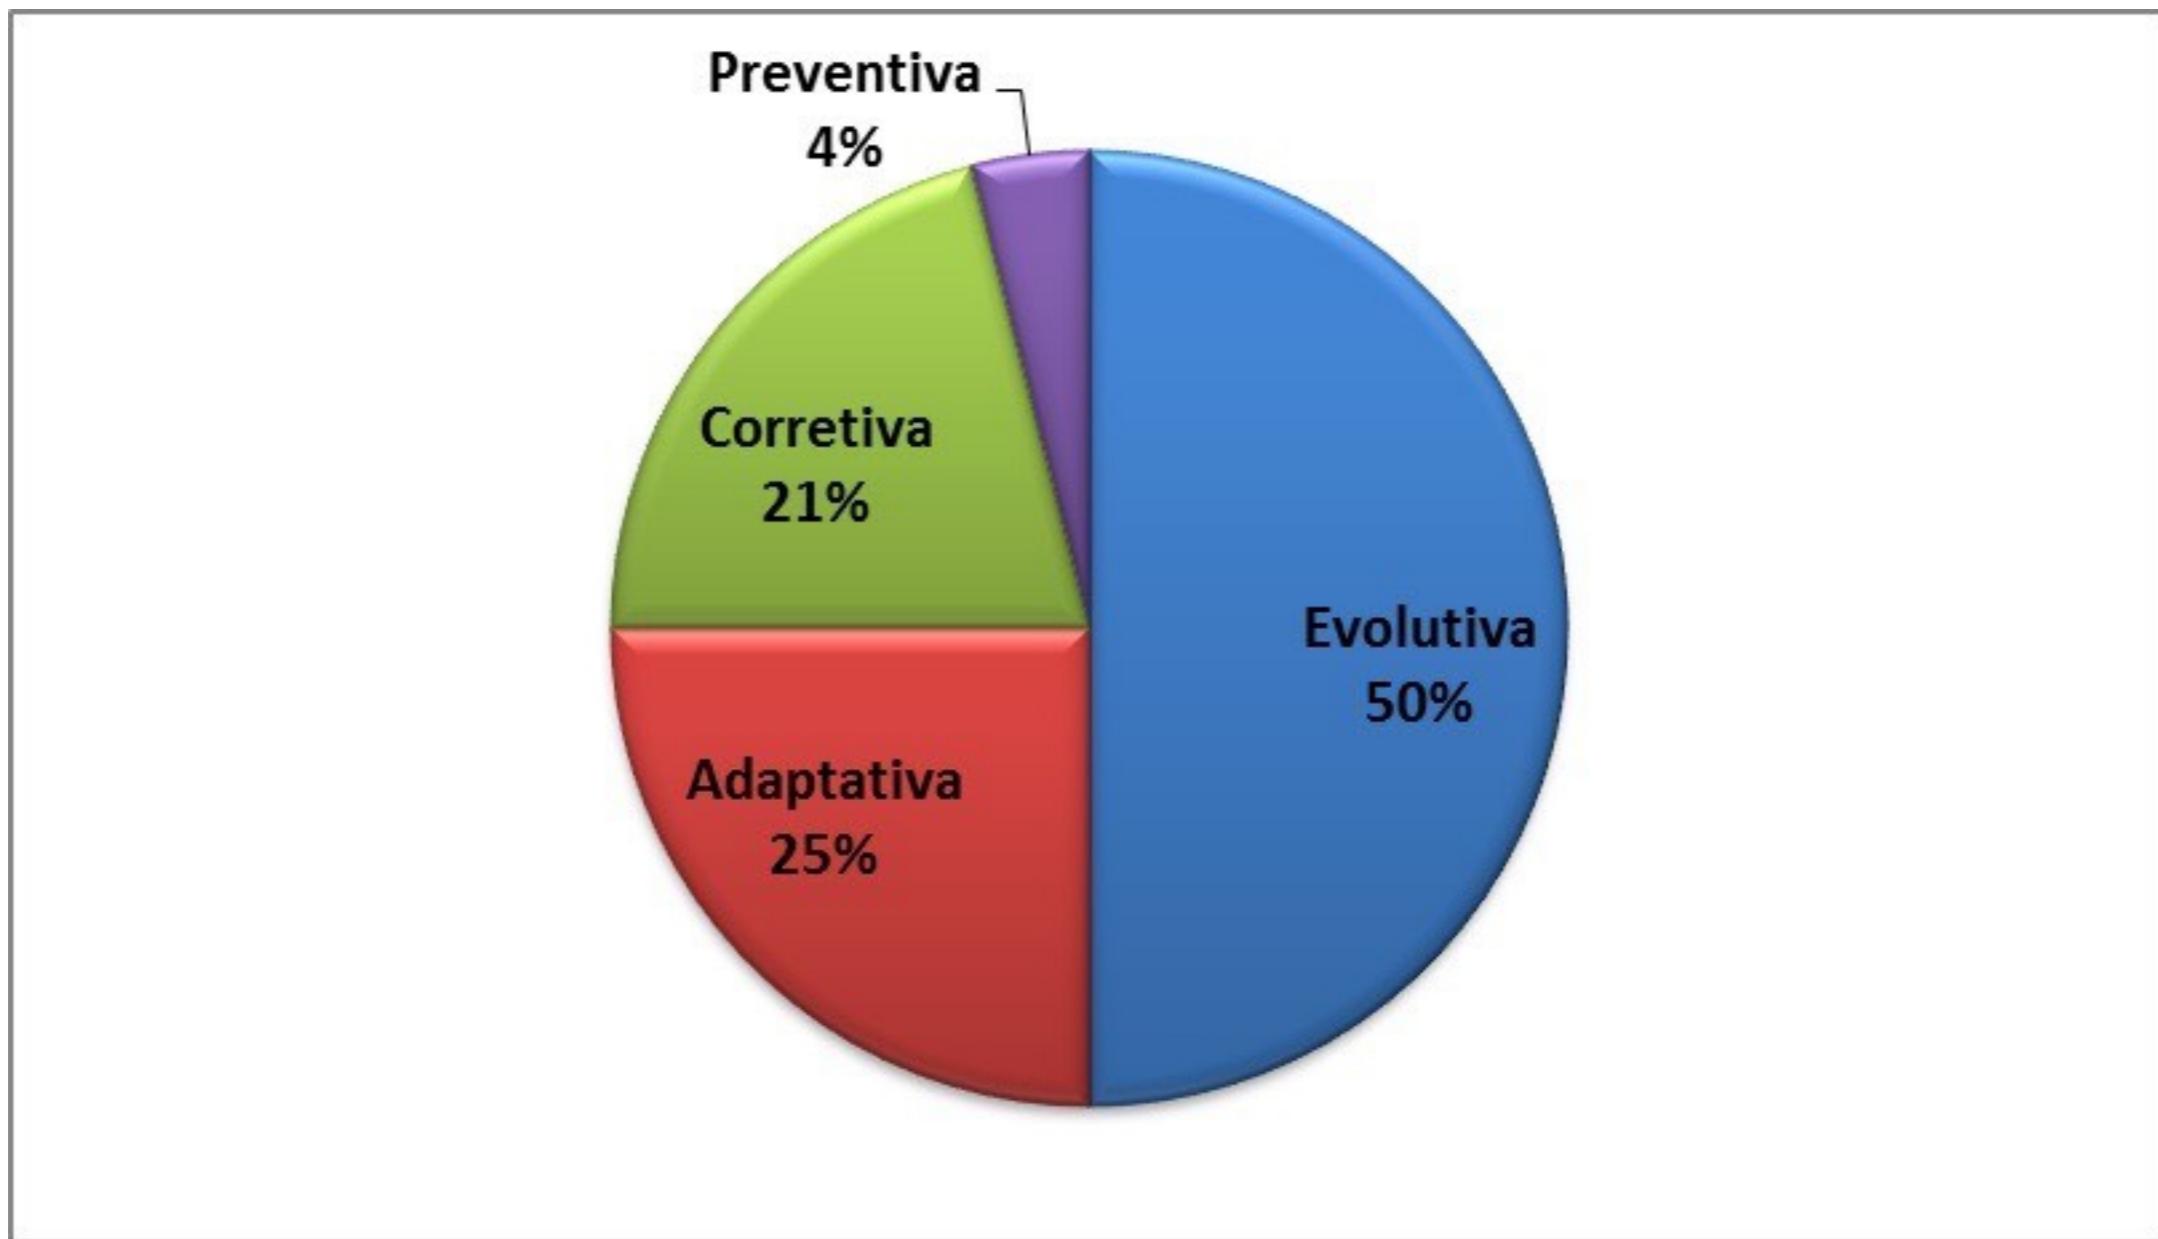
Mineração de Software: Introdução

Prof. André Hora
DCC/UFMG
2019.1

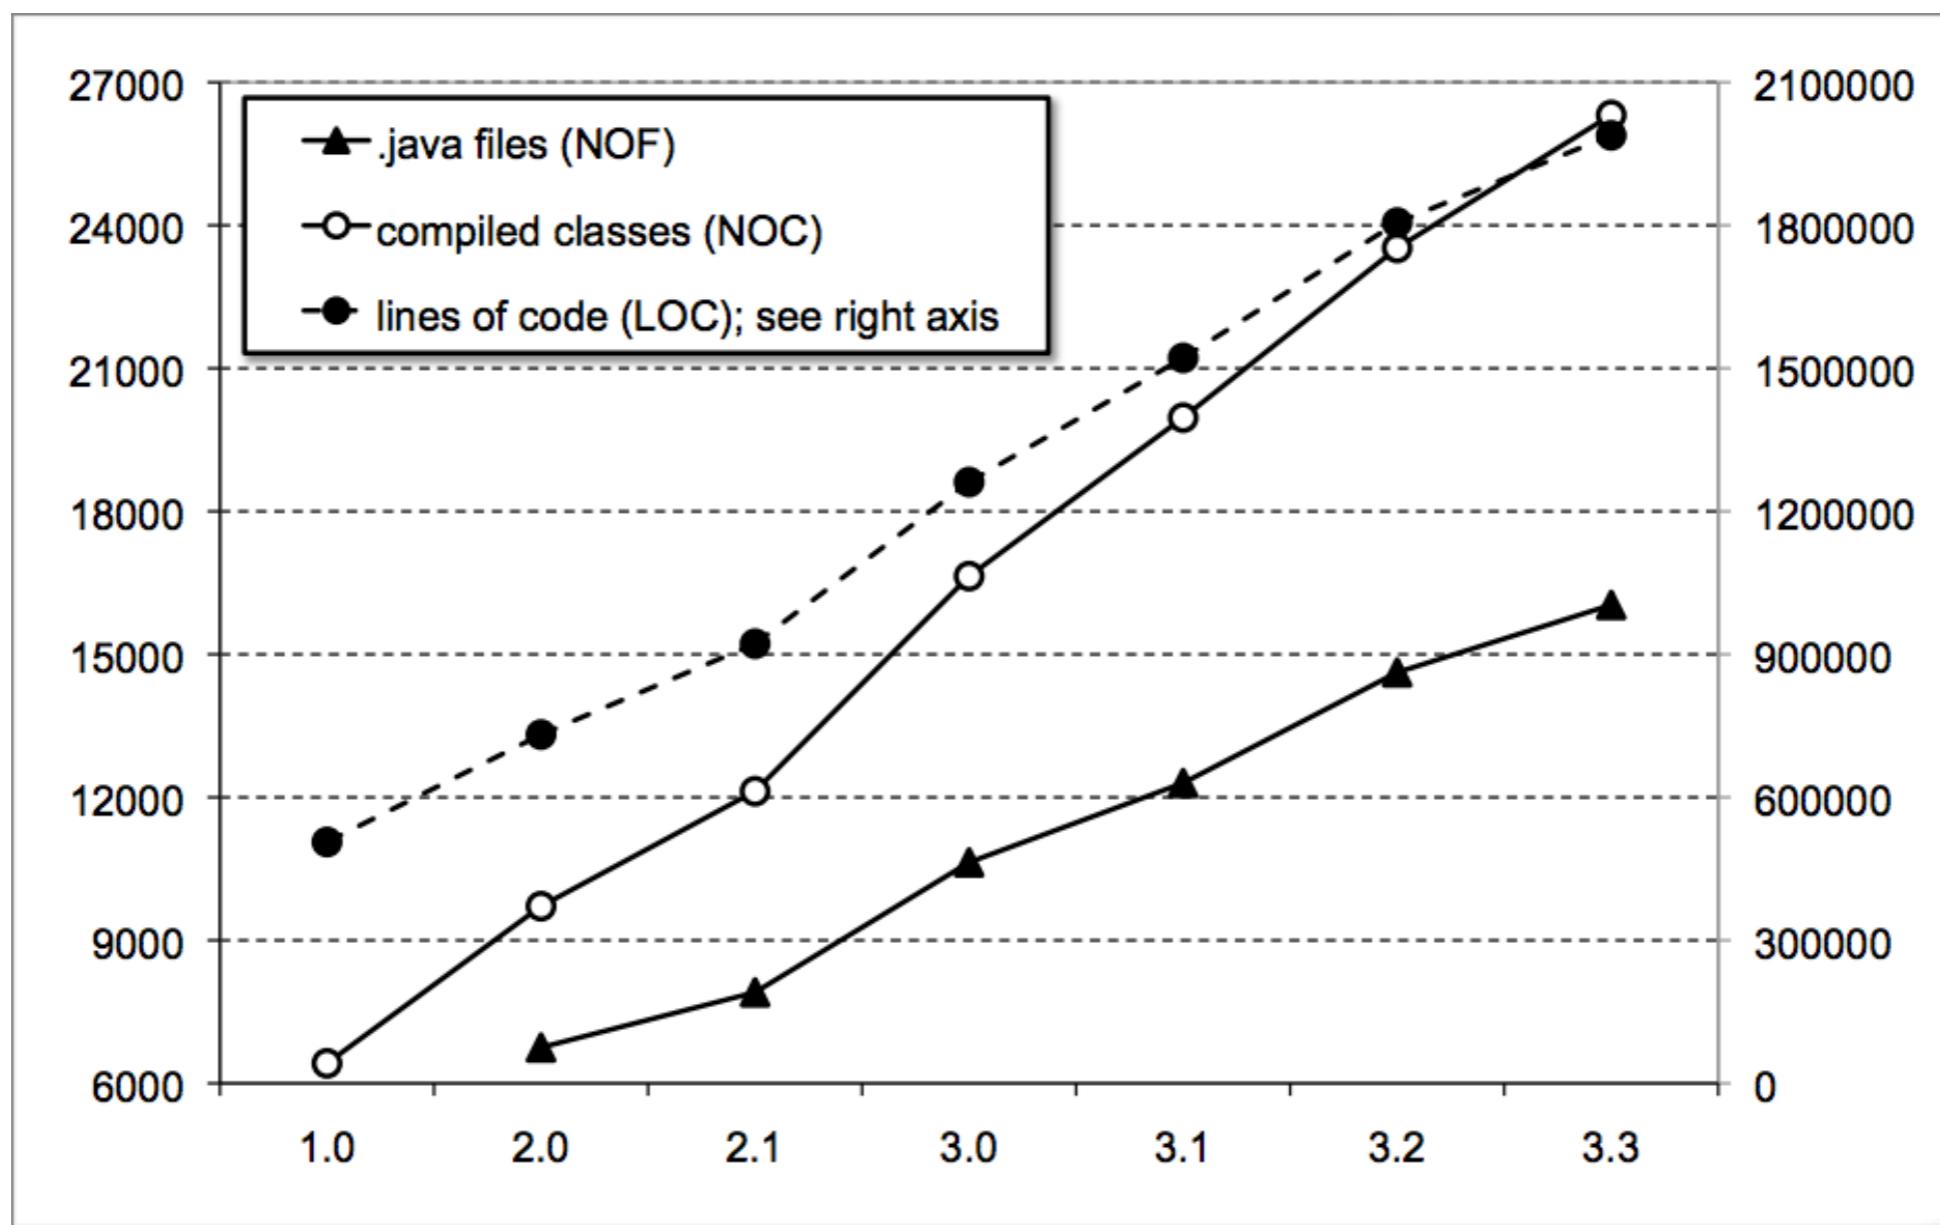
Evolução de Software

- Sistemas de software estão em constante evolução
- Novas funcionalidades, correção de bugs, refatoração de código
- Representa até **90%** dos custos de desenvolvimento

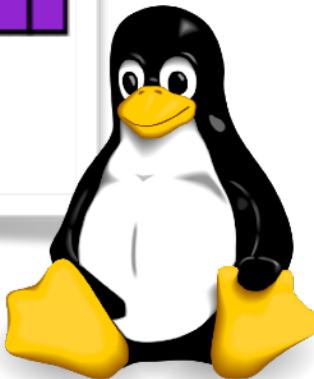
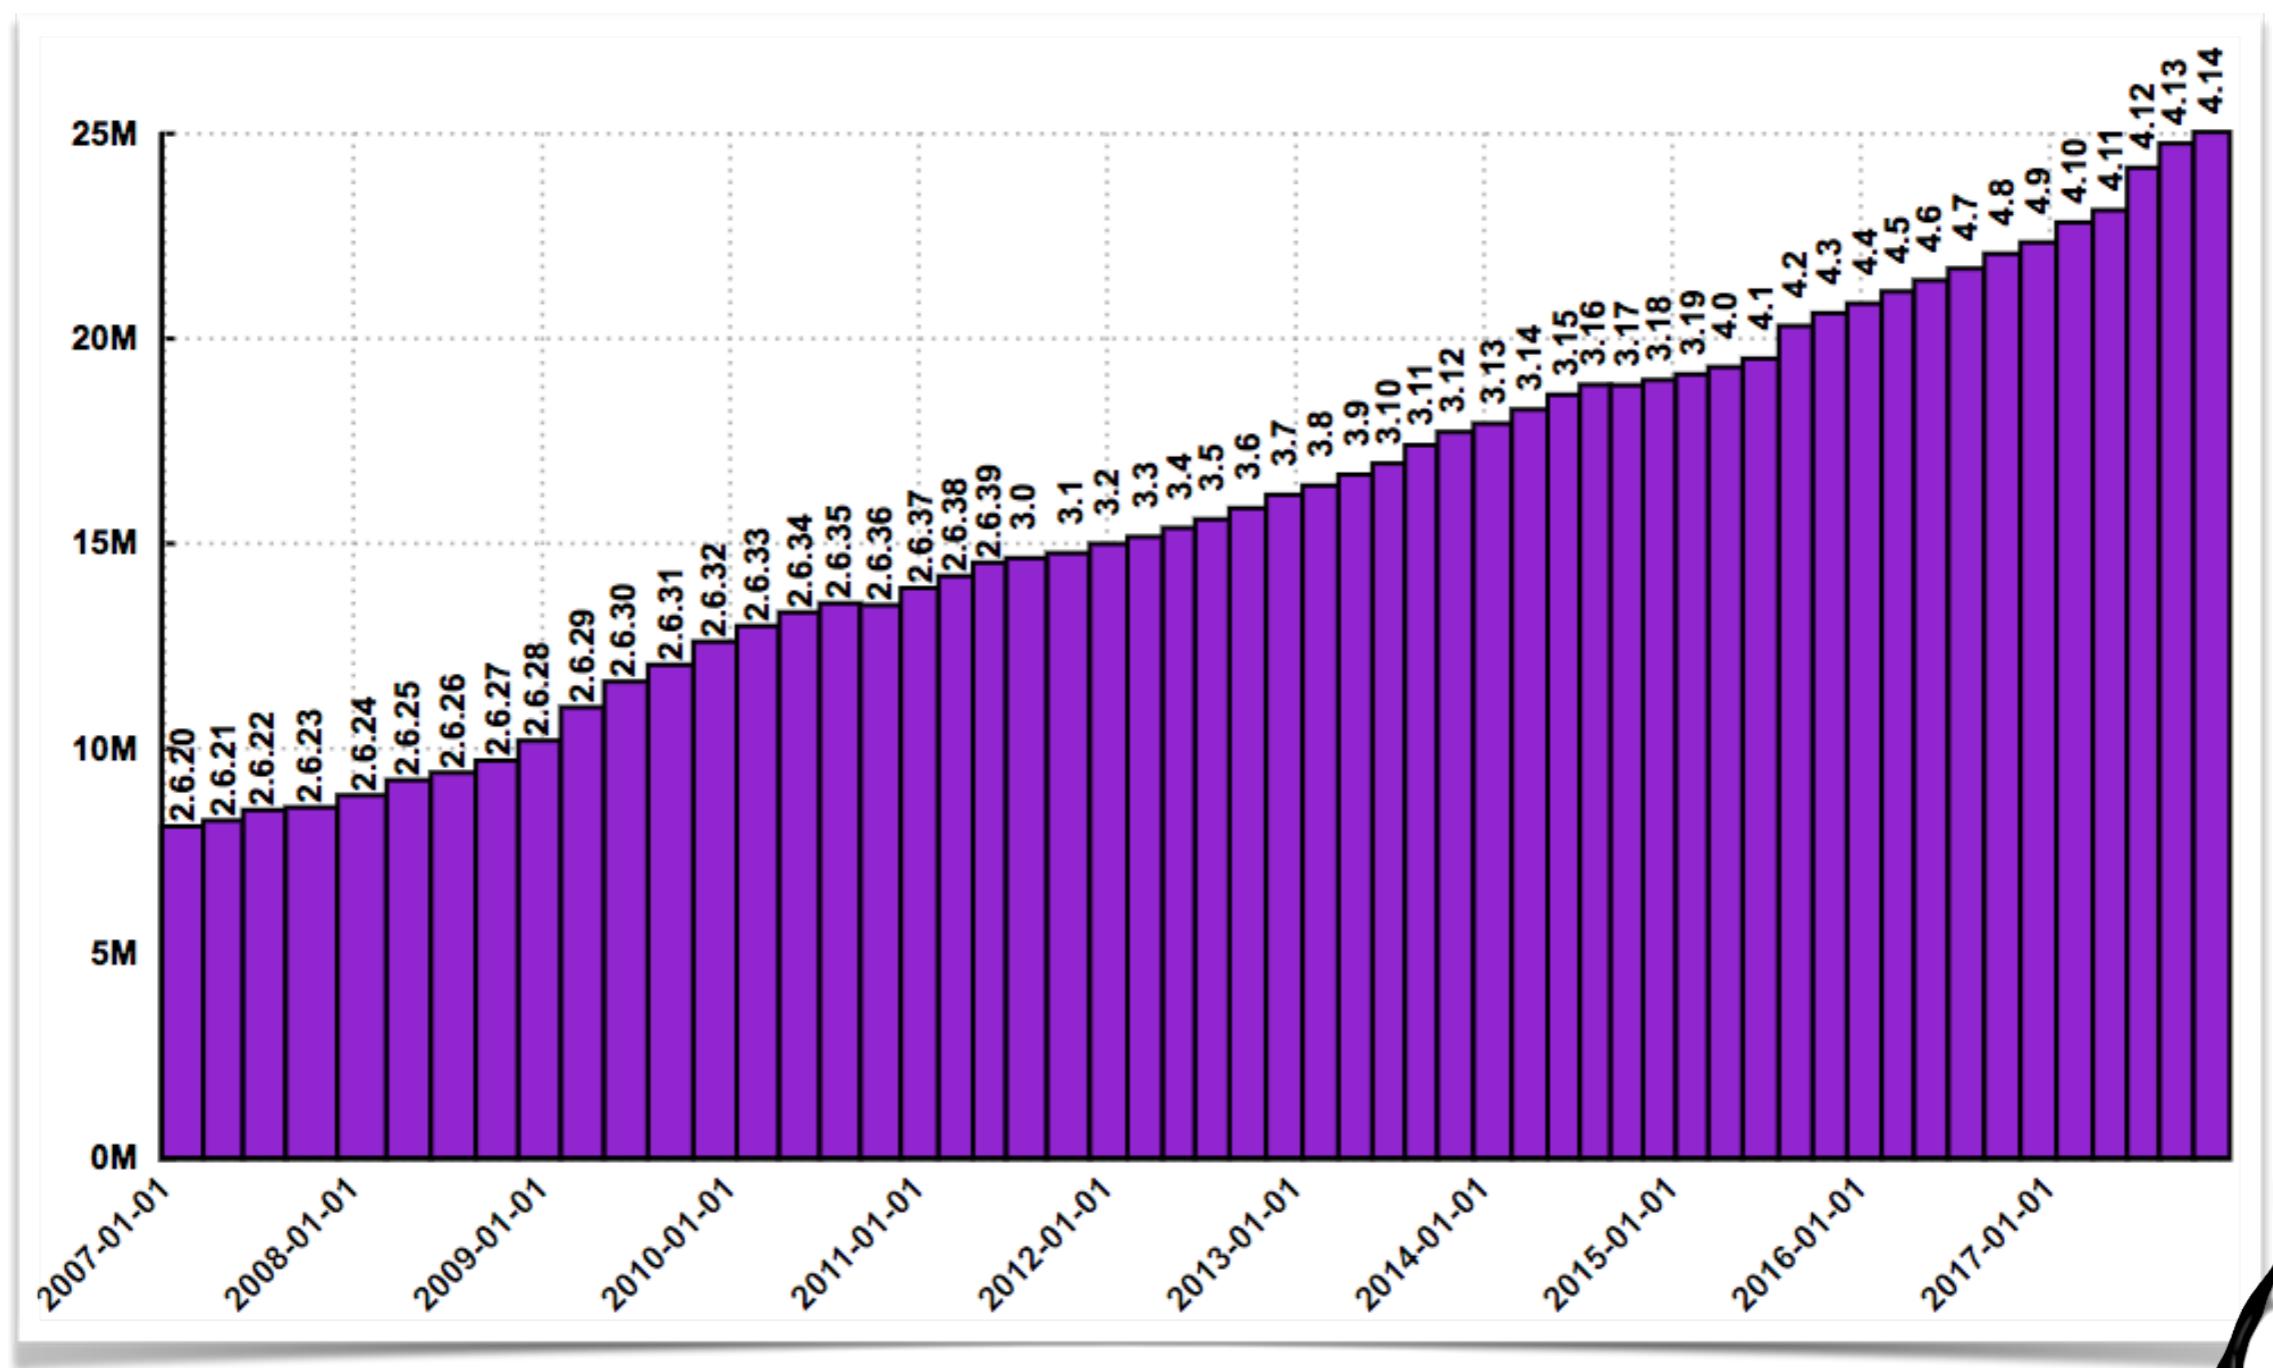
Proporção de Manutenção



Linhas de Código do Eclipse



Linhas de Código do Linux



Razões para Evolução de Software

**Se um sistema é utilizando, ele nunca está finalizado
pois precisa sempre evoluir para:**

Acomodar novas funcionalidades

Corrigir defeitos

Melhorar design

Comunicar com outros sistemas

Migrar de SO, BD, bibliotecas, etc

Adaptar a diferentes hardware

Adaptar a leis, regras de negócio, etc

Refatorar código

Leis de Lehman

- ↑ Mudança contínua
- ↑ Complexidade crescente
- ↑ Crescimento contínuo
- ↓ Declínio da qualidade

Mineração de Software



Mineração de Software

- Reppositórios de software contém uma fonte rica de informações do estado **atual** e do seu **passado**
- Utilizando técnicas voltadas para mineração de repositórios de software:
 - **Desenvolvedores:** adquirem maior conhecimento sobre como manter e evoluir sistemas de software
 - **Pesquisadores:** exploraram esses dados para entender e melhorar as práticas de desenvolvimento

Mineração de Software

- Ajuda a melhorar nossa compreensão sobre o desenvolvimento de software
- Pode contribuir para a criação de novos métodos e de uma nova geração de ferramentas na ES
- Exemplos:
 - Detecção de códigos candidatos a refatoração
 - Suporte à migração de bibliotecas de software
 - Sugestão de melhoria de modularidade
 - Suporte a melhoria de documentação, entre vários outros...

Engenharia de Software II

Table I.1. The 15 SWEBOK KAs
Software Requirements
Software Design
Software Construction
Software Testing
Software Maintenance
Software Configuration Management
Software Engineering Management
Software Engineering Process
Software Engineering Models and Methods
Software Quality
Software Engineering Professional Practice
Software Engineering Economics
Computing Foundations
Mathematical Foundations
Engineering Foundations



**Mineração
de
Software**

Quais artefatos podem ser analisados?

Quais artefatos podem ser analisados?



Dados
(ex, código,
documentação)

Quais artefatos podem ser analisados?

Dados
(ex, código,
documentação)

Metadados
(ex, autor, avaliações
de um app,
downloads)

Quais artefatos podem ser analisados?

Dados
(ex, código,
documentação)

Histórico de
Versões
(ex, releases,
commits)

Metadados
(ex, autor, avaliações
de um app,
downloads)

Dados (eg, código, documentação)

```

for i in people.data.users:
    response = client.api.statuses.user_timeline.get(screen_name=i.screen_name)
    print 'Got', len(response.data), 'tweets from', i.screen_name
    if len(response.data) != 0:
        ldate = response.data[0]['created_at']
        ldate2 = datetime.strptime(ldate, '%a %b %d %H:%M:%S +0000 %Y')
        today = datetime.now()
        howlong = (today-ldate2).days
        if howlong < daywindow:
            print i.screen_name, 'has tweeted in the past', daywindow,
            totaltweets += len(response.data)
            for j in response.data:
                if j.entities.urls:
                    for k in j.entities.urls:
                        newurl = k['expanded_url']
                        urlset.add((newurl, j.user.screen_name))
        else:
            print i.screen_name, 'has not tweeted in the past', daywind
    
```

```

/**
 * Simple HelloButton() method.
 * @version 1.0
 * @author john doe <doe.j@example.com>
 */
HelloButton()
{
    JButton hello = new JButton( "Hello, world" );
    hello.addActionListener( new HelloBtnList
    // use the JFrame type until support for t
    // new component is finished
    JFrame frame = new JFrame( "Hello Button" );
    Container pane = frame.getContentPane();
    pane.add( hello );
    frame.pack();
    frame.show();           // display the frame
}

```

PACKAGE CLASS USE TREE DEPRECATED INDEX HELP
PREV CLASS NEXT CLASS FRAMES NO FRAMES
SUMMARY NESTED FIELD CONSTR METHOD DETAIL FIELD CONSTR METHOD
com.baeldung.javadoc
Class SuperHero
java.lang.Object
com.baeldung.javadoc.Person
com.baeldung.javadoc.SuperHero

public class SuperHero
extends Person

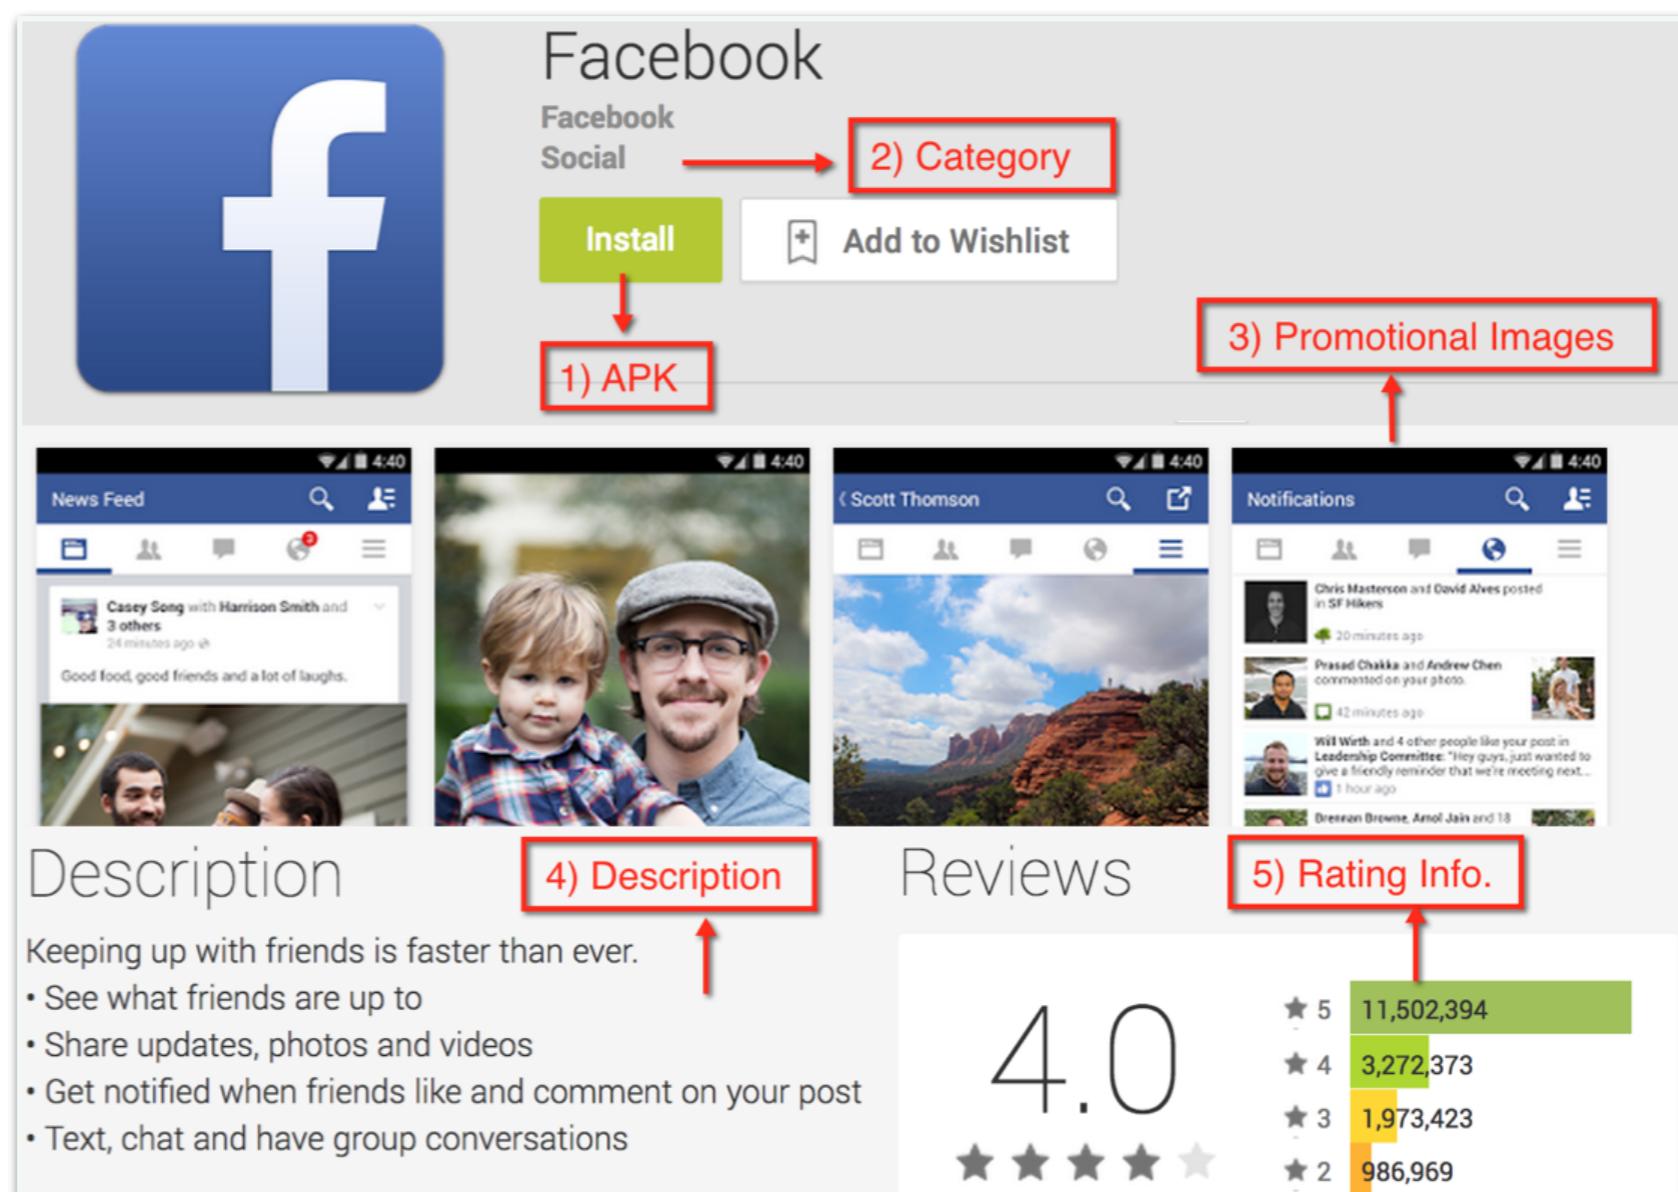
Hero is the main entity we will be using to ...

Author:
Captain America

Field Summary

Fields	Modifier and Type	Field and Description
	private int	defense
	private int	health
	private java.lang.String	heroName The public name of a hero that is common knowledge
	private java.lang.String	uniquePower

Metadados (eg, autor, avaliações de um app, downloads)



Metadados (eg, autor, avaliações de um app, downloads)

Change sort order of shard stats in CCR test

This commit changes the sort order of shard stats that are collected in CCR retention lease integration tests. This change is done so that primaries appear first in sort order.



master (#140) 

 jasonitedor committed 2 days ago  1 parent 91f69fc commit 14510a933a0b1ecd38033ecc21f2621eba45698a

2 x-pack/plugin/CCR/src/test/java/org/elasticsearch/xpack/CCRRetentionLeaseIT.java 

```
@@ -568,7 +568,7 @@ public void testUnfollowFailsToRemoveRetentionLeases() throws Exception {  
 568     return Arrays.stream(stats.getShards())  
 569         .sorted((s, t) -> {  
 570             if (s.getShardRouting().shardId().id() ==  
 t.getShardRouting().shardId().id()) {  
 571 -                 return  
 Boolean.compare(s.getShardRouting().primary(),  
 t.getShardRouting().primary());  
 572             } else {  
 573                 return  
 Integer.compare(s.getShardRouting().shardId().id(),  
 t.getShardRouting().shardId().id());  
 574         }  
 568     return Arrays.stream(stats.getShards())  
 569 +         .sorted((s, t) -> {  
 570             if (s.getShardRouting().shardId().id() ==  
 t.getShardRouting().shardId().id()) {  
 571 +                 return  
 Boolean.compare(s.getShardRouting().primary(),  
 t.getShardRouting().primary());  
 572             } else {  
 573                 return  
 Integer.compare(s.getShardRouting().shardId().id(),  
 t.getShardRouting().shardId().id());  
 574         }  
 575 }
```

Histórico de Versões (eg, releases, commits)



```
$ git log
commit ca82a6dff817ec66f44342007202690a93763949
Author: Scott Chacon <schacon@gee-mail.com>
Date:   Mon Mar 17 21:52:11 2008 -0700

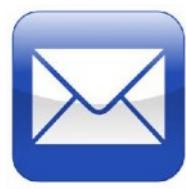
    changed the version number

commit 085bb3bcb608e1e8451d4b2432f8ecbe6306e7e7
Author: Scott Chacon <schacon@gee-mail.com>
Date:   Sat Mar 15 16:40:33 2008 -0700

    removed unnecessary test

commit a11bef06a3f659402fe7563abf99ad00de2209e6
Author: Scott Chacon <schacon@gee-mail.com>
Date:   Sat Mar 15 10:31:28 2008 -0700

    first commit
```

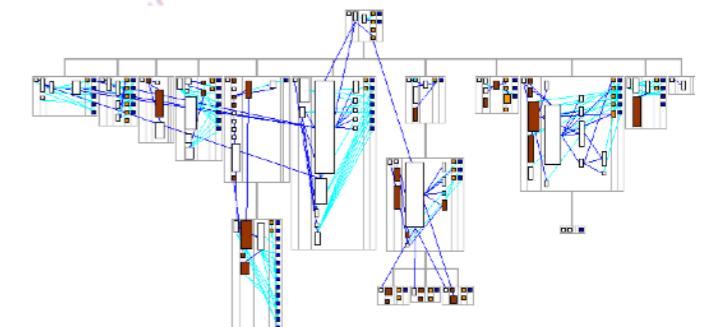
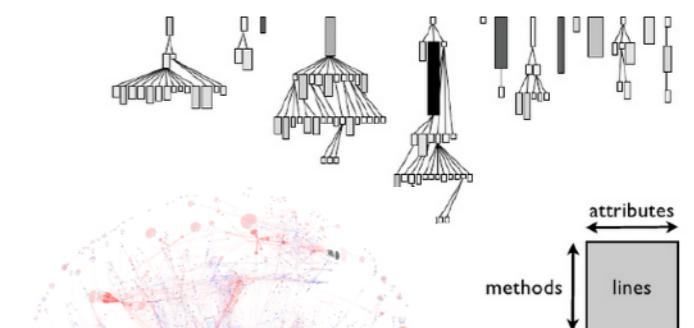
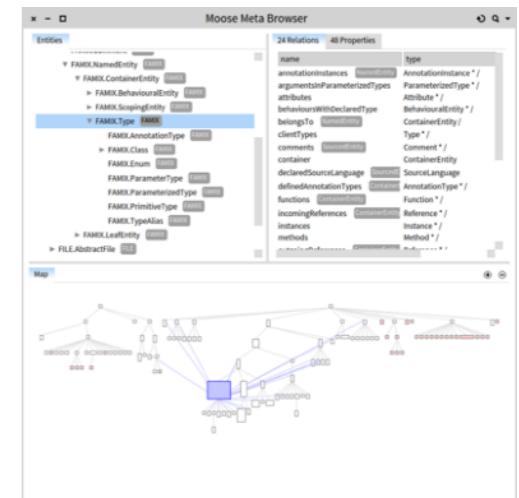
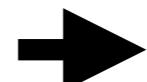


Bugzilla

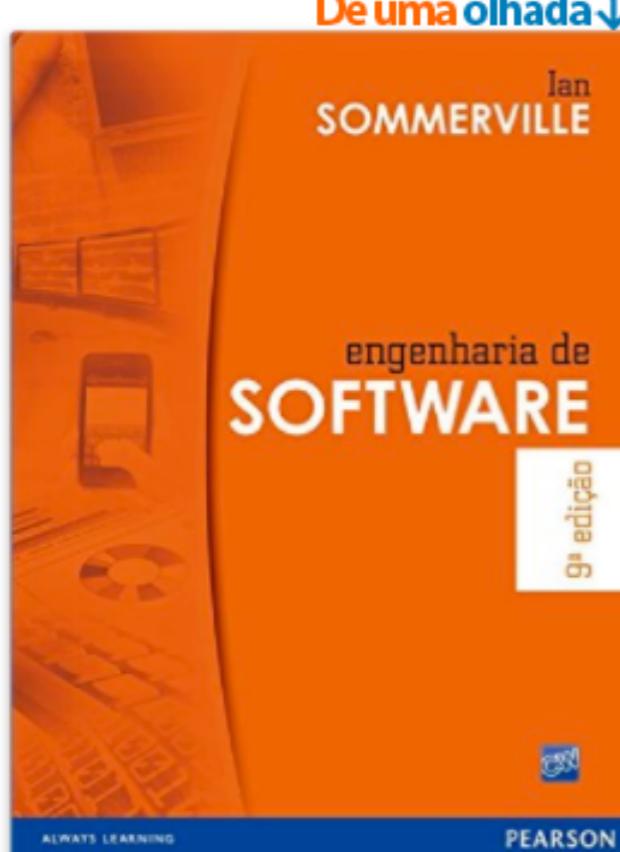
```
1 (function() {  
2  
3     window.utils = window.utils || {};  
4  
5     /*  
6      Random Number Generator.  
7  
8      Pretty awesome explanation here:  
9      http://stackoverflow.com/a/1527820  
*/  
10    utils.randomNumber = function(min, max) {  
11        return Math.floor(Math.random() * (max - min + 1)) + min;  
12    };  
13}
```



Mineração
de
Software



Exemplos



[Ver todas as 2 imagens](#)

Dê uma olhada↓

Engenharia de Software (Português) Capa Comum – 15 jun 2011

por [Ian Sommerville](#) (Autor)

8 avaliações de clientes

[Ver todos os 3 formatos e edições](#)

eBook Kindle

R\$ 149,10

Capa Comum

R\$ 156,95

[Leia com nossos apps gratuitos](#)

11 Novo(s) a partir de R\$ 156,95

Em até 5x R\$ 31,39 sem juros [Calculadora de prestações](#)

Entrega para o CEP 01319-900 na Sexta-feira, 1 de Setembro, se você finalizar o pedido em **23 horas e 20 minutos**

Leia

Enquanto
Enviamos

Leia Enquanto Enviamos

Compre e comece a ler a amostra digital deste livro enquanto espera ele chegar. Saiba mais [aqui](#).

No intuito de atender às necessidades de alunos e professores dos cursos de ciência da computação, engenharia de computação e sistema de informação, esta nona edição de Engenharia de software teve seu conteúdo reestruturado e totalmente atualizado. A obra conta agora com novos capítulos que focalizam o desenvolvimento ágil de software e os sistemas embarcados, além de trazer novas abordagens sobre engenharia dirigida a modelos, desenvolvimento open source, modelo Swiss Cheese de Reason,

[Leia mais](#)

Frequentemente comprados juntos



Preço total: **R\$ 308,93**

[Adicionar ambos ao carrinho](#)

Este item: Engenharia de Software por Ian Sommerville Capa comum **R\$ 156,95**

Engenharia de Software. Uma Abordagem Profissional por Roger S. Pressman Capa comum **R\$ 151,98**



[Ver todas as 2 imagens](#)

Dê uma olhada

Engenharia de Software (Português) Capa Comum – 15 jun 2011

por Ian Sommerville (Autor)

★★★★★ 8 avaliações de clientes

[Ver todos os 3 formatos e edições](#)

eBook Kindle

R\$ 149,10

Capa Comum

R\$ 156,95

[Leia com nossos apps gratuitos](#)

11 Novo(s) a partir de R\$ 156,95

Em até 5x R\$ 31,39 sem juros [Calculadora de prestações](#)

Entrega para o CEP 01319-900 na Sexta-feira, 1 de Setembro, se você finalizar o pedido em **23 horas e 20 minutos**

Leia

Enquanto
Enviamos

Leia Enquanto Enviamos

Compre e comece a ler a amostra digital deste livro enquanto espera ele chegar. Saiba mais [aqui](#).

No intuito de atender às necessidades de alunos e professores dos cursos de ciência da computação, engenharia de computação e sistema de informação, esta nona edição de Engenharia de software teve seu conteúdo reestruturado e totalmente atualizado. A obra conta agora com novos capítulos que focalizam o desenvolvimento ágil de software e os sistemas embarcados, além de trazer novas abordagens sobre engenharia dirigida a modelos, desenvolvimento open source, modelo Swiss Cheese de Reason,

[Leia mais](#)

Frequentemente comprados juntos



Preço total: **R\$ 308,93**

[Adicionar ambos ao carrinho](#)

Este item: Engenharia de Software por Ian Sommerville Capa comum **R\$ 156,95**

Engenharia de Software. Uma Abordagem Profissional por Roger S. Pressman Capa comum **R\$ 151,98**

Data mining



[Ver todas as 2 imagens](#)

Frequentemente comprados juntos



Preço total: **R\$ 308,93**

[Adicionar ambos ao carrinho](#)

Este item: Engenharia de Software por Ian Sommerville Capa comum **R\$ 156,95**

Engenharia de Software. Uma Abordagem Profissional por Roger S. Pressman Capa comum **R\$ 151,98**

Engenharia de Software (Português) Capa Comum – 15 jun 2011

por [Ian Sommerville](#) (Autor)

8 avaliações de clientes

[Ver todos os 3 formatos e edições](#)

eBook Kindle

R\$ 149,10

Capa Comum

R\$ 156,95

[Leia com nossos apps gratuitos](#)

11 Novo(s) a partir de R\$ 156,95

Em até 5x R\$ 31,39 sem juros [Calculadora de prestações](#)

Entrega para o CEP 01319-900 na Sexta-feira, 1 de Setembro, se você finalizar o pedido em **23 horas e 20 minutos**

Leia

[Enquanto Enviamos](#)

Leia Enquanto Enviamos

Compre e comece a ler a amostra digital deste livro enquanto espera ele chegar. Saiba mais [aqui](#).

No intuito de atender às necessidades de alunos e professores dos cursos de ciência da computação, engenharia de computação e sistema de informação, esta nona edição de Engenharia de software teve seu conteúdo reestruturado e totalmente atualizado. A obra conta agora com novos capítulos que focalizam o desenvolvimento ágil de software e os sistemas embarcados, além de trazer novas abordagens sobre engenharia dirigida a modelos, desenvolvimento open source, modelo Swiss Cheese de Reason,

[Leia mais](#)

Data mining

Sommerville → Pressman

Podemos utilizar essa técnica no contexto de software?

Reglas de uso de APIs

```
234             String lock = unlock.getText().toString();
235             lock = lock.substring(0, lock.length() - 1);
236             unlock.setText(lock);
...
327             mHomeKeyLocker.lock(this);
328
329         }
330
331         unlock.setFilters(new InputFilter[] { new InputFilter.LengthFilter(12)
332     });
}
```

before → after

lock() → unlock()

openFile() → closeFile()

openStream() → closeStream()

Regras de uso de APIs

```
234         String lock = unlock.getText().toString();
235         lock = lock.substring(0, lock.length() - 1);
236         unlock.setText(lock);
...
327         mHomeKeyLocker.lock(this);
328     }
329
330
331     unlock.setFilters(new InputFilter[] { new InputFilter.LengthFilter(12)
332 });
}
```

before → after

lock() → unlock()

openFile() → closeFile()

openStream() → closeStream()

Qual a aplicação real?

Regras de uso de APIs

```
234         String lock = unlock.getText().toString();
235         lock = lock.substring(0, lock.length() - 1);
236         unlock.setText(lock);
...
327         mHomeKeyLocker.lock(this);
328     }
329
330
331     unlock.setFilters(new InputFilter[] { new InputFilter.LengthFilter(12)
332 });
}
```

before → after

lock() → unlock()

openFile() → closeFile()

openStream() → closeStream()

Qual a aplicação real?

- Melhorar documentação
- Detectar bugs

Regras de evolução de APIs

```
package org.jboss.as.modcluster;

-import junit.framework.Assert;
import org.jboss.as.controller.PathAddress;
import org.jboss.dmr.ModelNode;
import org.jboss.dmr.ModelType;
+import org.junit.Assert;
import org.junit.Test;

    public Collection<FormulaData> getChildren() {
-    List<FormulaData> result = new ArrayList<FormulaData>();
+    List<FormulaData> result = Lists.newArrayList();
        for (DecoratorContext childContext : decoratorContext.getChildren()) {
            result.add(new DefaultFormulaData(childContext));
        }
    }
```

old → new

`ArrayList()` → `List.newArrayList()`

`Vector` → `ArrayList`

`junit.framework.Assert` → `org.junit.Assert`

Regras de evolução de APIs

```
package org.jboss.as.modcluster;

-import junit.framework.Assert;
import org.jboss.as.controller.PathAddress;
import org.jboss.dmr.ModelNode;
import org.jboss.dmr.ModelType;
+import org.junit.Assert;
import org.junit.Test;

    public Collection<FormulaData> getChildren() {
-    List<FormulaData> result = new ArrayList<FormulaData>();
+    List<FormulaData> result = Lists.newArrayList();
        for (DecoratorContext childContext : decoratorContext.getChildren()) {
            result.add(new DefaultFormulaData(childContext));
        }
    }
```

Qual a aplicação real?

old → new

`ArrayList()` → `List.newArrayList()`

`Vector` → `ArrayList`

`junit.framework.Assert` → `org.junit.Assert`

Regras de evolução de APIs

```
package org.jboss.as.modcluster;

-import junit.framework.Assert;
import org.jboss.as.controller.PathAddress;
import org.jboss.dmr.ModelNode;
import org.jboss.dmr.ModelType;
+import org.junit.Assert;
import org.junit.Test;

    public Collection<FormulaData> getChildren() {
-    List<FormulaData> result = new ArrayList<FormulaData>();
+    List<FormulaData> result = Lists.newArrayList();
        for (DecoratorContext childContext : decoratorContext.getChildren()) {
            result.add(new DefaultFormulaData(childContext));
        }
    }
```

Qual a aplicação real?

old → new

ArrayList() → List.newArrayList()

- Melhorar documentação
- Migrar bibliotecas

Vector → ArrayList

junit.framework.Assert → org.junit.Assert

Mais exemplos?

Pesquisas em MSR

Mineração de Software

- Reppositórios de software contém uma fonte rica de informações do estado **atual** e do seu **passado**
- Utilizando técnicas voltadas para mineração de repositórios de software:
 - **Desenvolvedores:** adquirem maior conhecimento sobre como manter e evoluir sistemas de software
 - **Pesquisadores:** exploraram esses dados para entender e melhorar as práticas de desenvolvimento

Conferências

- **MSR**: International Conference on Mining Software Repositories
- **ICSE**: ACM/IEEE International Conference on Software Engineering
- **ICSME**: IEEE International Conference on Software Maintenance and Evolution
- **SANER**: IEEE International Conference on Software Analysis, Evolution and Reengineering
- **VEM**: Workshop on Software Visualization, Evolution and Maintenance

[Attending](#) ▾[Program](#) ▾[Tracks](#) ▾[Organization](#) ▾ [Search](#)[Series](#) ▾[Sign in](#)[Sign up](#)

[Attending](#) ▾[Program](#) ▾[Tracks](#) ▾[Organization](#) ▾ [Search](#)[Series](#) ▾[Sign in](#)[Sign up](#)

Gothenburg

35th IEEE International Conference on Software Maintenance and Evolution

Cleveland, OH USA

September 30th - October 4th 2019



[Attending ▾](#)[Program ▾](#)[Tracks ▾](#)[Organization ▾](#) [Search](#)[Series ▾](#)[Sign in](#)[Sign up](#)

International Conference on Maintenance and Evolution

Cleveland, OH USA
September 30th - October 4th 2019



[Attending](#) ▾[Program](#) ▾[Tracks](#) ▾[Organization](#) ▾[Search](#)[Series](#) ▾[Sign in](#)[Sign up](#)

Attending ▾

Program ▾

Tracks ▾

Organization ▾

Search

Series ▾

Sign in

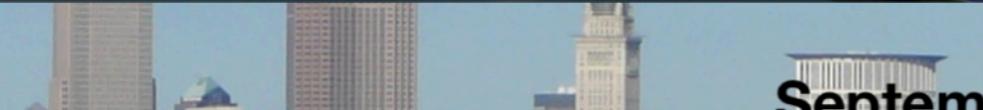
Sign up



Hangzhou, China
February 24-27, 2019



**International Conference on
Maintenance and Evolution**



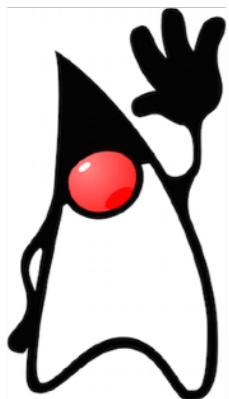
VI Workshop on Software Visualization, Evolution and Maintenance

**Cleveland, OH USA
September 30th - October 4th 2019**

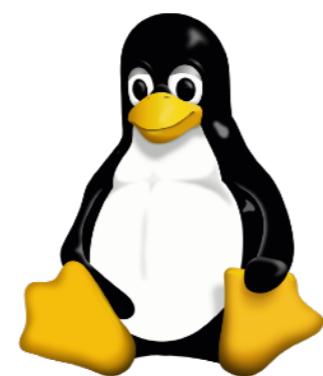
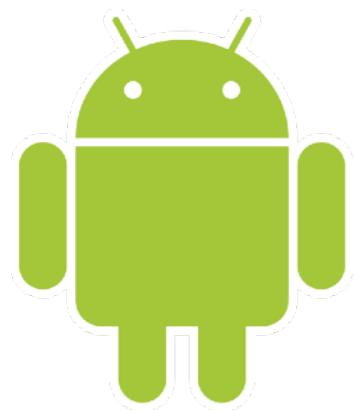
VEM 2018

September, 2018
São Carlos, Brazil





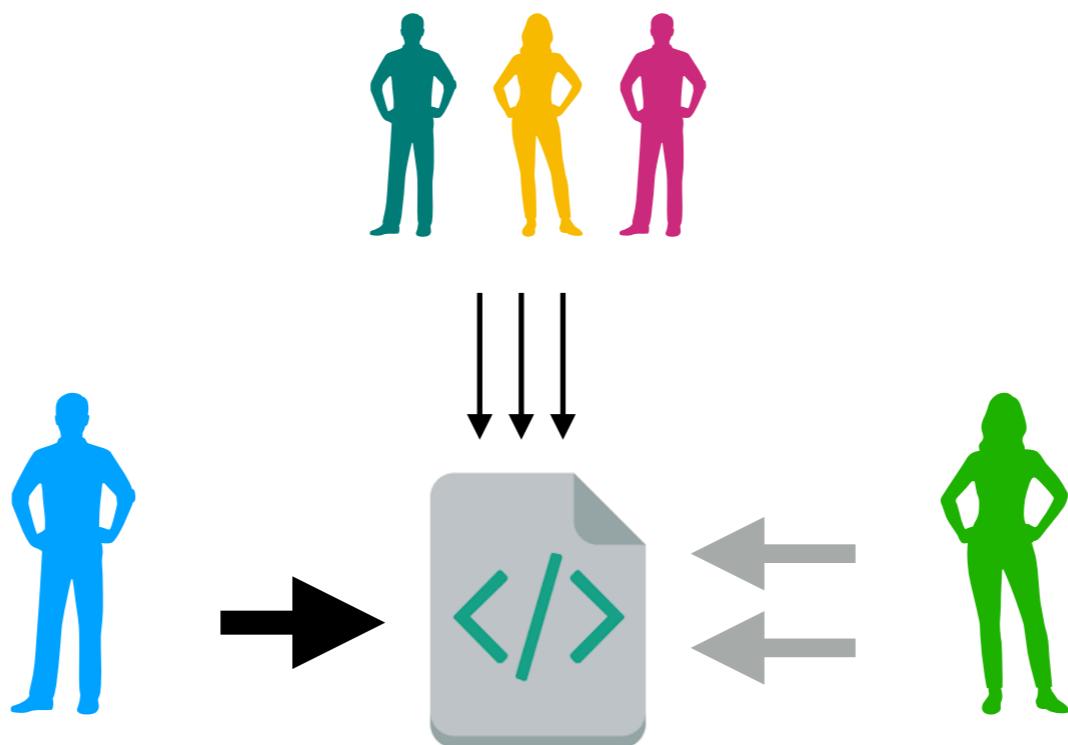
Jenkins



Travis CI



Pesquisas (exemplos)



Who Can Maintain this Code?

Assessing the Effectiveness of Repository-Mining Techniques for Identifying Software Maintainers

Guilherme Avelino^{1,2}, Leonardo Passos³, Fabio Petrillo⁴, Marco Túlio Valente¹

¹Federal University of Minas Gerais, Brazil

²Federal University of Piauí, Brazil

³Quantstamp Technologies, Canada

⁴Polytechnique Montréal, Canada

{gaa, mtov}@dcc.ufmg.br, leo@quantstamp.com, fabio@petrillo.com

Abstract

In large and complex systems, identifying developers capable of maintaining a piece of code is an important but challenging task. In this context, repository-mining techniques can help by providing some level of automation. Still, whether such techniques effectively

```
while (!stack.isEmpty()) {
```

```
    State current = stack.pop();
```

```
    //  
    //  
    //  
    int theDepth = current.getDepth();  
    if(theDepth > maxDepthSearched)  
        maxDepthSearched = theDepth;
```

*Commented-out code:
Código “removido”
através de comentário*

```
    if (current.isGoal(maze)) {  
        // TODO update the maze if a solution found  
        //while(current != null){  
            while(current.getParent() != null){  
                if(!current.isGoal(maze))  
                    maze.setOneSquare(current.getSquare(), '.');  
                current = current.getParent();  
            }  
        return true;  
    }
```

Minerando Código Comentado

Lucas Grijó¹, Andre Hora¹

¹ Faculdade de Computação (FACOM)

Universidade Federal de Mato Grosso do Sul (UFMS)

rksgrijo@gmail.com, hora@facom.ufms.br

Abstract. As software evolves, programmers often comment code snippets as a mean of removing them (a practice known as comment-out code). However, this bad practice may cause problems to software maintainers, such as readability reduction, distraction, and waste of time. In this context, this paper presents an empirical study to evaluate the presence of commented-out code on open source systems. We analyze 100 relevant software systems and 300 releases. As a result, we detect a rate of 4,17% commented-out code. However, we verify that

```
package org.jboss.as.modcluster;
-import junit.framework.Assert;
import org.jboss.as.controller.PathAddress;
import org.jboss.dmr.ModelNode;
import org.jboss.dmr.ModelType;
+import org.junit.Assert;
import org.junit.Test;

    public Collection<FormulaData> getChildren() {
-    List<FormulaData> result = new ArrayList<FormulaData>();
+    List<FormulaData> result = Lists.newArrayList();
        for (DecoratorContext childContext : decoratorContext.getChildren()) {
            result.add(new DefaultFormulaData(childContext));
        }
    }
```

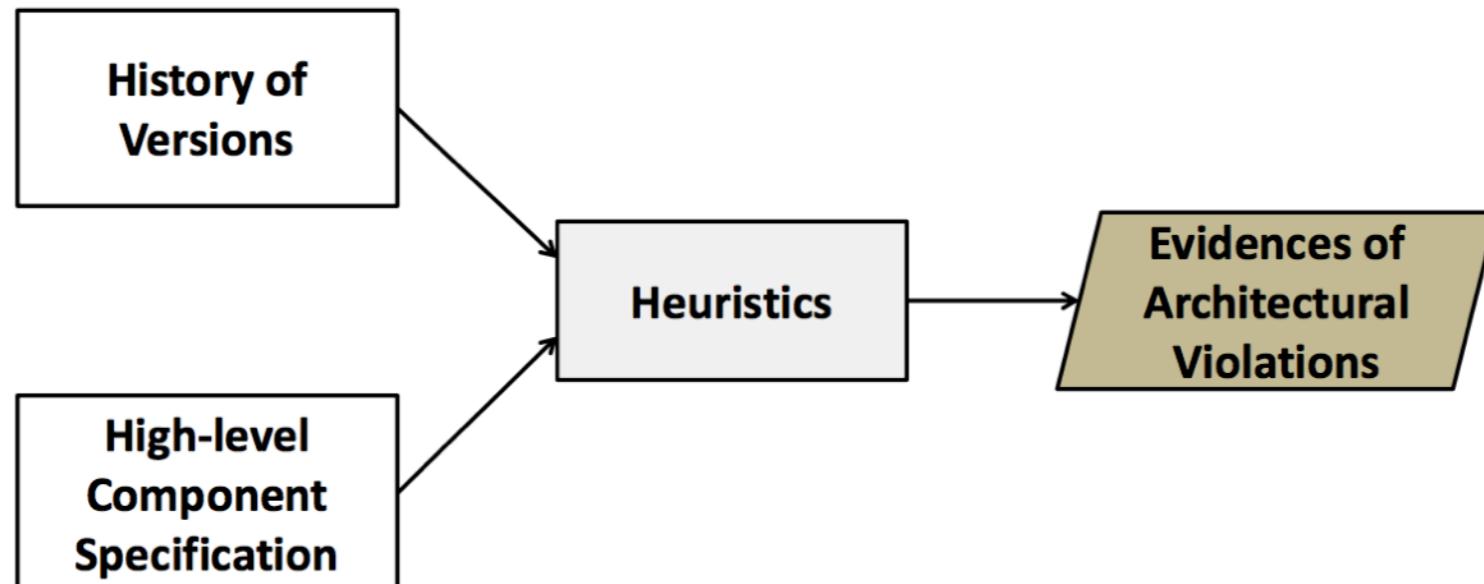
Minerando Mensagens de Depreciação Faltantes em APIs: Um Estudo de Caso no Ecossistema Android

Pedro Henrique de Moraes¹, Caroline Lima¹, Andre Hora¹

¹ Faculdade de Computação (FACOM)
Universidade Federal de Mato Grosso do Sul (UFMS)

pedhmoraes@gmail.com, carollimaxp@gmail.com, hora@facom.ufms.br

Abstract. *To facilitate library migration, developers should deprecate APIs before any change that may impact on client systems. However, the literature reports that APIs are commonly deprecated without any message to support their clients. Therefore, in this paper, we propose an approach to recommend missing replacement messages in deprecated APIs. We assess the history version of the Android framework and 100 client systems. We found that some replacement messages can be detected by mining the clients, but alternatives solutions are common. Finally, based on our learning, we suggest improvements to the design of an API deprecation migration tool.*

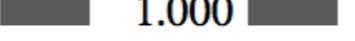
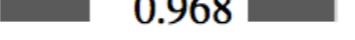
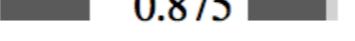
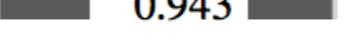
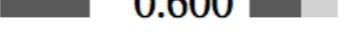
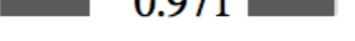
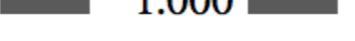
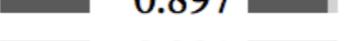
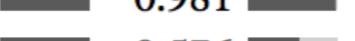
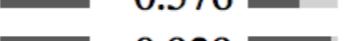
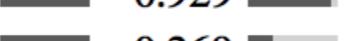
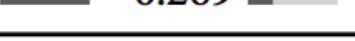


Mining Architectural Violations from Version History

**Cristiano Maffort · Marco Túlio Valente ·
Ricardo Terra · Mariza Bigonha · Nicolas
Anquetil · André Hora**

Received: date / Accepted: date

Abstract Software architecture conformance is a key software quality control activity that aims to reveal the progressive gap normally observed between concrete and planned software architectures. However, formally specifying an architecture can be difficult, as it must be done by an expert of the system having a high level understanding of it. In this paper, we present a lightweighted approach for architecture conformance based on a combination of static and historical source code analysis. The proposed approach relies on four heuristics for detecting absences (something expected was

RDiff		
Ref. Type	Precision	Recall
Rename Type	1.000	
Move Type	1.000	
Extract Superclass	1.000	
Rename Method	1.000	
Pull Up Method	1.000	
Push Down Method	1.000	
Move Method	1.000	
Extract Method	1.000	
Inline Method	1.000	
Pull Up Field	1.000	
Push Down Field	1.000	
Move Field	1.000	

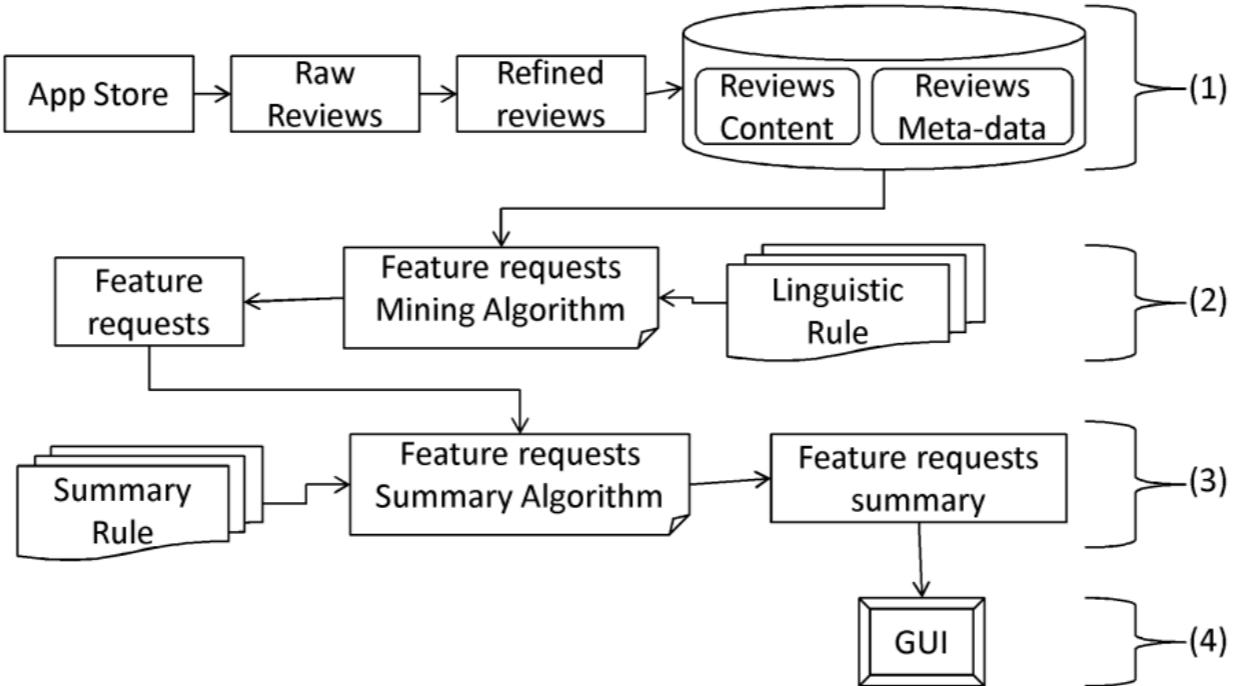
RefDiff: Detecting Refactorings in Version Histories

Danilo Silva*, Marco Túlio Valente†
 Department of Computer Science
 Universidade Federal de Minas Gerais
 Belo Horizonte, Brazil
 Email: *danilos@dcc.ufmg.br, †mtov@dcc.ufmg.br

Abstract—Refactoring is a well-known technique that is widely adopted by software engineers to improve the design and enable the evolution of a system. Knowing which refactoring operations were applied in a code change is a valuable information to understand software evolution, adapt software components, merge code changes, and other applications. In this paper, we present RefDiff, an automated approach that identifies refactorings performed between two code revisions in a git repository. RefDiff employs a combination of heuristics based on static analysis and code similarity to detect 13 well-known refactoring types. In an evaluation using an oracle of 418 known refactoring operations,

the refactorings that were applied to an API, we can replay them on the client code automatically.

Although there are approaches capable of detecting refactorings automatically, there are still some issues that hinder their application. Specifically, the precision and recall of such approaches still need improvements. In this paper, we try to fill this gap by proposing RefDiff, an automated approach that identifies refactorings performed in the version history of a system. RefDiff employs a combination of heuristics



Retrieving and Analyzing Mobile Apps Feature Requests from Online Reviews

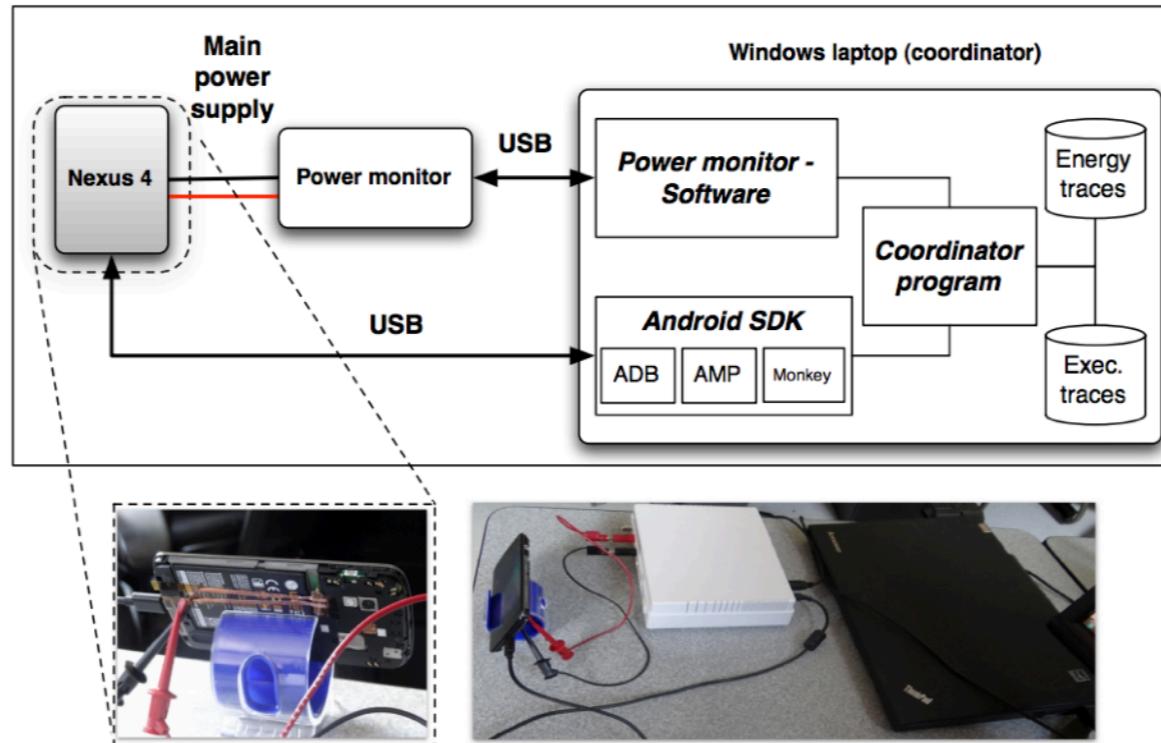
Claudia Iacob, Rachel Harrison

Department of Computing and Communication Technologies
Oxford Brookes University
Oxford, United Kingdom
[{iacob, rachel.harrison}@brookes.ac.uk](mailto:{iacob,rachel.harrison}@brookes.ac.uk)

Abstract—Mobile app reviews are valuable repositories of ideas coming directly from app users. Such ideas span various topics, and in this paper we show that 23.3% of them represent feature requests, i.e. comments through which users either suggest new features for an app or express preferences for the re-design of already existing features of an app. One of the challenges app developers face when trying to make use of such feedback is the massive amount of available reviews. This makes it difficult to identify specific topics and recurring trends across reviews. Through this work, we aim to support such processes by designing MARA (Mobile App Review Analyzer), a prototype for automatic retrieval of mobile app feature requests from online reviews. The design of the prototype is a) informed by an investigation of the ways users express feature requests through reviews, b) developed around a set of pre-defined linguistic rules, and c) evaluated on a large sample of online reviews. The results

Facebook app alone has almost 6 million reviews. Moreover, the most popular non free apps on the Google App store get an average of 45.5 reviews per day. Second, reviews have a style of their own; they are usually short, unstructured and seldom obey grammar and punctuation rules. Moreover, users express opinions in unconventional ways, often using unconventional syntax or sarcasm to vent their feelings after purchasing an app. Even though humans can usually interpret such reviews, it is challenging to automate such interpretations [6].

We first considered a sample of reviews for analysis and we aimed to identify how much of the users' feedback consists of feature requests and how these feature requests are expressed through reviews. Such findings informed the design of a prototype system able to identify and retrieve feature



Mining Energy-Greedy API Usage Patterns in Android Apps: An Empirical Study

Mario Linares-Vásquez¹, Gabriele Bavota², Carlos Bernal-Cárdenas³

Rocco Oliveto⁴, Massimiliano Di Penta², Denys Poshyvanyk¹

¹The College of William and Mary, Williamsburg, VA, USA

²University of Sannio, Benevento, Italy

³Universidad Nacional de Colombia, Bogotá, Colombia

⁴University of Molise, Pesche (IS), Italy

mlinarev@cs.wm.edu, gbavota@unisannio.it, cebernalc@unal.edu.co,

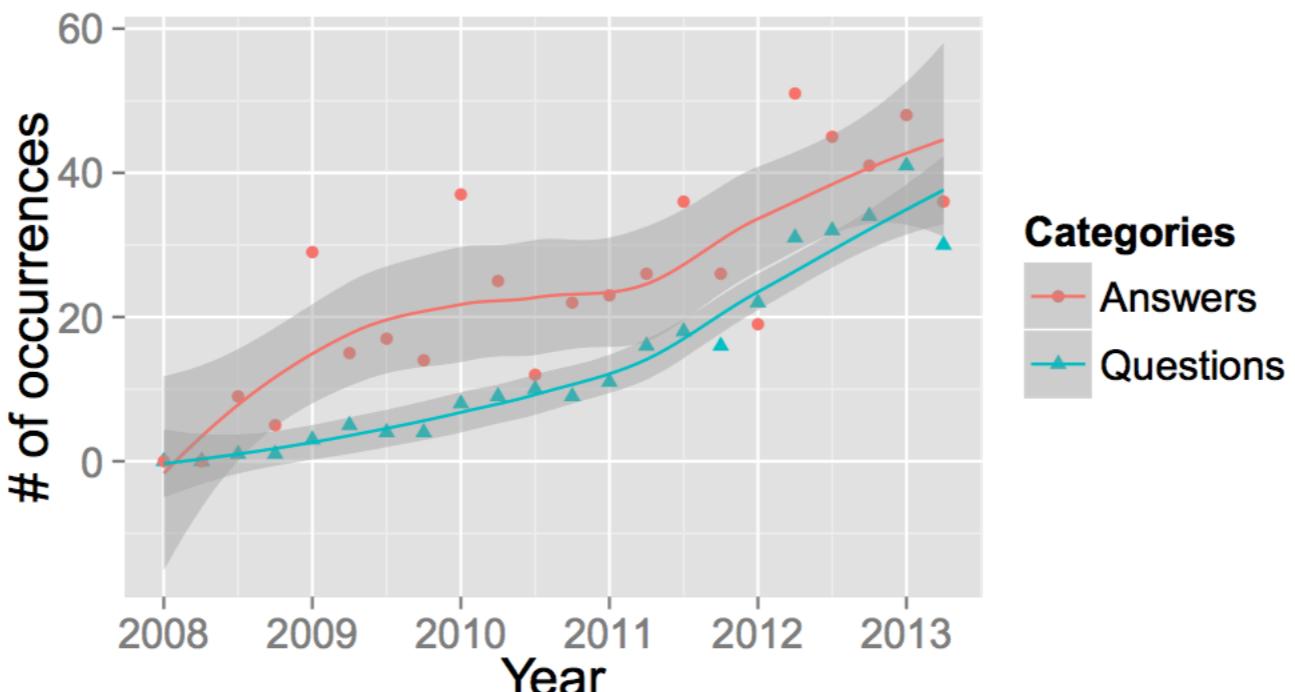
rocco.oliveto@unimol.it, dipenta@unisannio.it, denys@cs.wm.edu

ABSTRACT

Energy consumption of mobile applications is nowadays a hot topic, given the widespread use of mobile devices. The high demand for features and improved user experience, given the available powerful hardware, tend to increase the apps' energy consumption. However, excessive energy consumption in mobile apps could also be a consequence of energy greedy hardware, bad programming practices, or particular API usage patterns. We present the largest to date quantitative and qualitative empirical investigation into the categories of API calls and usage patterns that—in the context of the

1. INTRODUCTION

In recent years, we are observing rapid evolution of mobile devices in terms of available hardware, operating systems, and, as a consequence of that, the growing lists of features that mobile applications' (apps) users demand. These modern apps have virtually the same features as their equivalent desktop applications. For instance, many top video games for mobile devices provide similar levels of user experience as compared to those console analogs. Such evident step-ahead has, however, a price to be paid. Nowadays, multi-core processors, high-performance Graphical Processing Units



Mining Questions about Software Energy Consumption

Gustavo Pinto^{†‡}, Fernando Castor[†], Yu David Liu[‡]

[†]Federal University of Pernambuco
Recife, PE, Brazil
{ghlp, castor}@cin.ufpe.br

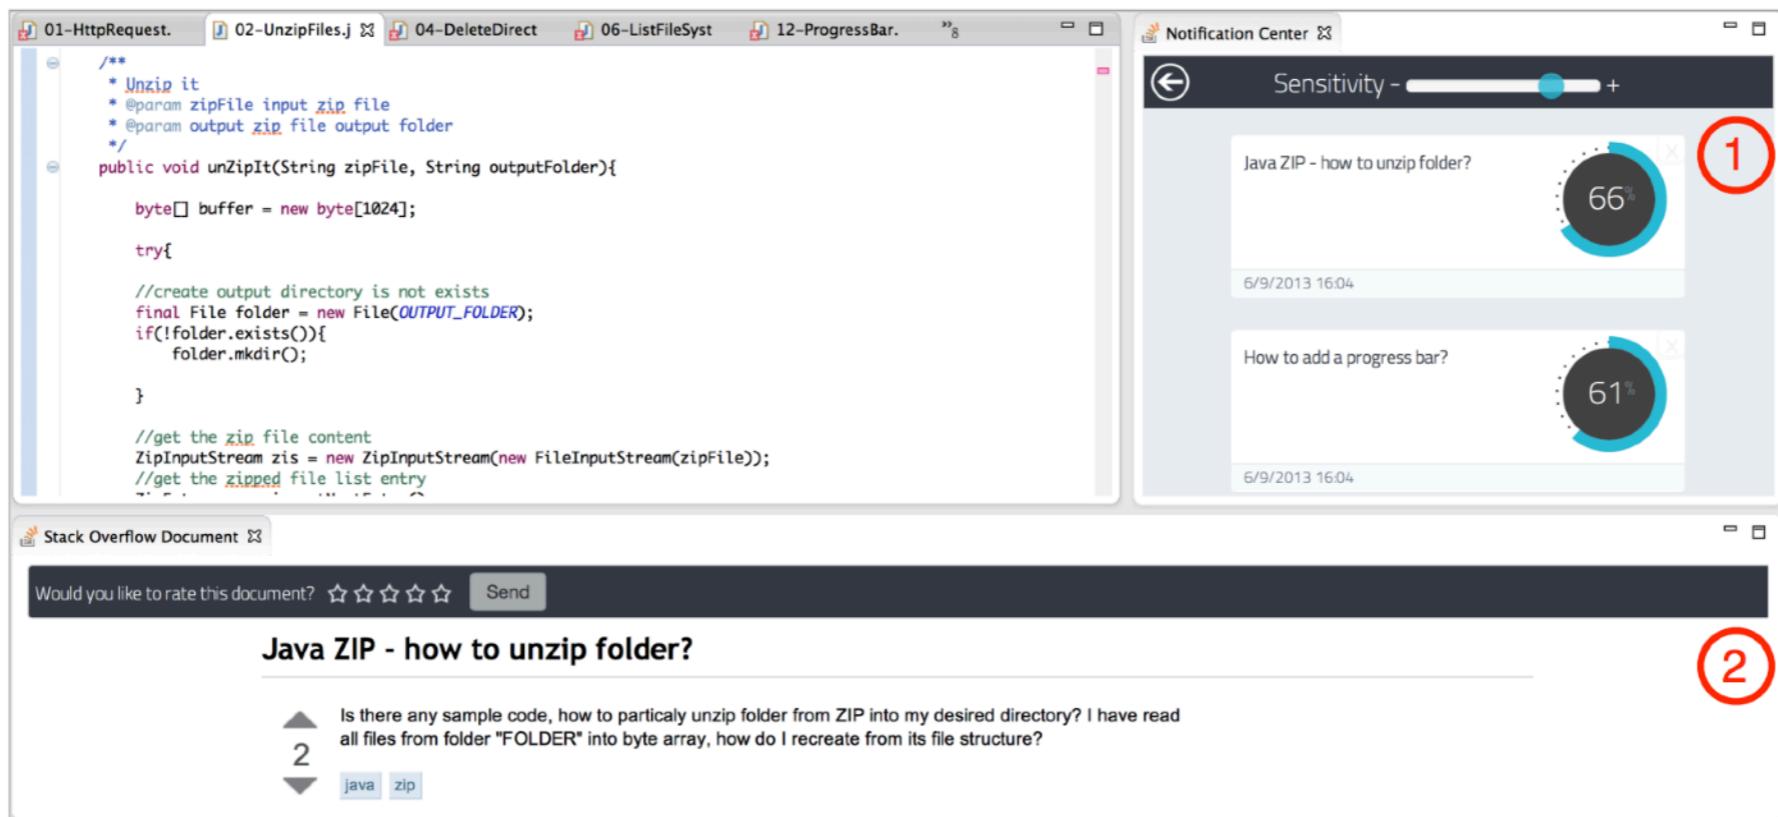
[‡]SUNY Binghamton
Binghamton, NY 13902, USA
davidL@cs.binghamton.edu

ABSTRACT

A growing number of software solutions have been proposed to address application-level energy consumption problems in the last few years. However, little is known about how much software developers are concerned about energy consumption, what aspects of energy consumption they consider important, and what solutions they have in mind for improving energy efficiency. In this paper we present the first empirical study on understanding the views of application programmers on software energy consumption problems. Using STACKOVERFLOW as our primary data source, we analyze a carefully curated sample of more than 300 questions and 550 answers from more than 800 users. With this data, we

1. INTRODUCTION

Nowadays, thanks to the rapid proliferation of mobile phones, tablets, and unwired devices in general, energy efficiency is becoming a key software design consideration where the energy consumption is closely related to battery lifetime. It is also of increasing interest in the non-mobile arena, such as data centers and desktop environments. Energy-efficient solutions are highly sought after across the compute stack, with more established results through innovations in hardware/architecture [2, 14, 28], operating systems [10, 19, 24], and runtime systems [8, 25, 31]. In recent years, there is a growing interest in studying energy consumption from higher layers of the compute stack and most of these stud-



Mining StackOverflow to Turn the IDE into a Self-Confident Programming Prompter

Luca Ponzanelli¹, Gabriele Bavota², Massimiliano Di Penta², Rocco Oliveto³, Michele Lanza¹

1: REVEAL @ Faculty of Informatics – University of Lugano, Switzerland

2: University of Sannio, Benevento, Italy

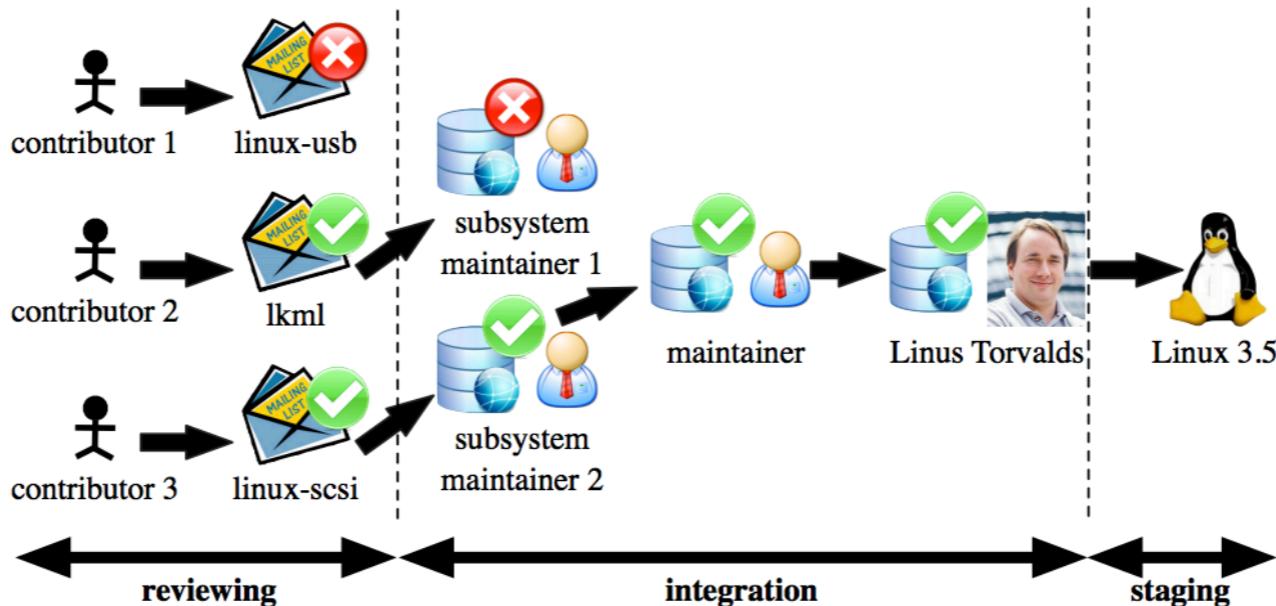
3: University of Molise, Pesche (IS), Italy

ABSTRACT

Developers often require knowledge beyond the one they possess, which often boils down to consulting sources of information like Application Programming Interfaces (API) documentation, forums, Q&A websites, *etc.* Knowing what to search for and how is non-trivial, and developers spend time and energy to formulate their problems as queries and to peruse and process the results.

We propose a novel approach that, given a context in the IDE,

problems, the main one being the absence of automation: Every time developers need to look for information, they interrupt their work flow, leave the IDE, and use a Web browser to perform and refine searches, and assess the results. Finally, they transfer the obtained knowledge to the problem context in the IDE. The information is retrieved from different sources, such as forums, mailing lists [2], blogs, Q&A websites, bug trackers [1], *etc.* A prominent example is Stack Overflow, popular among developers as a venue for sharing programming knowledge. Stack Overflow in 2010 it had



Will My Patch Make It? And How Fast?

Case Study on the Linux Kernel

Yujuan Jiang, Bram Adams
MCIS, Polytechnique Montréal, Canada
{jujuan.jiang,bram.adams}@polymtl.ca

Daniel M. German
University of Victoria, Canada
dmg@uvic.ca

Abstract—The Linux kernel follows an extremely distributed reviewing and integration process supported by 130 developer mailing lists and a hierarchy of dozens of Git repositories for version control. Since not every patch can make it and of those that do, some patches require a lot more reviewing and integration effort than others, developers, reviewers and integrators need support for estimating which patches are worthwhile to spend effort on and which ones do not stand a chance. This paper cross-links and analyzes eight years of patch reviews from the kernel mailing lists and committed patches from the Git repository to understand which patches are accepted and how long it takes those patches to get to the end user. We found that 33% of the patches makes it into a Linux release, and that most of them need 3 to 6 months for this. Furthermore, that patches developed by more experienced developers are more easily accepted and faster reviewed and integrated. Additionally, reviewing time is impacted by the number of authors, the number of commits, the size of the commit message, and the number of files.

ones, whose changes did not make it after putting months of work into them (e.g., [12], [13]). Even major projects like the Google Android mobile platform have problems getting their Linux kernel modifications integrated into the official kernel version [11]. Yet, determining up front whether a patch will make it, and how long it will take, is a grey area. Research on code reviews has shown how small patches [6], [4] sent by experienced developers [2] are more likely to be accepted by the reviewers, but it is not clear if these characteristics play the same role during the actual integration of the patch with other patches. Similarly, the impact of these characteristics on the time it takes to get a patch into a release is unclear.

This paper studies the relation of patch characteristics with (1) the probability of acceptance into an official release and